

Lab 11

Basic Image Processing
Fall 2018

Disclaimer

Accurate and important details regarding Lab11 materials (both theory & practice) were presented on the whiteboard of the classroom, upcoming slides are only sketchy summary of those informations.

SIFT

Scale Invariant Feature Transform

Feature: “is a piece of information which is relevant for solving the computational task related to a certain application” (Wikipedia)

Descriptor: characterization of a feature or keypoint


(these are superficial definitions, please see Lecture slides for more accurate shaping)

Main steps:

1. scale-space extrema detection (*build the pyramid*)
2. keypoint localization (*local extrema selection + rejection of weak candidates*)
3. orientation assignment
4. keypoint description


Why is it good for us?

(example) aim at: have a match between the “appropriate” parts of the images



Definitions

image pyramid:



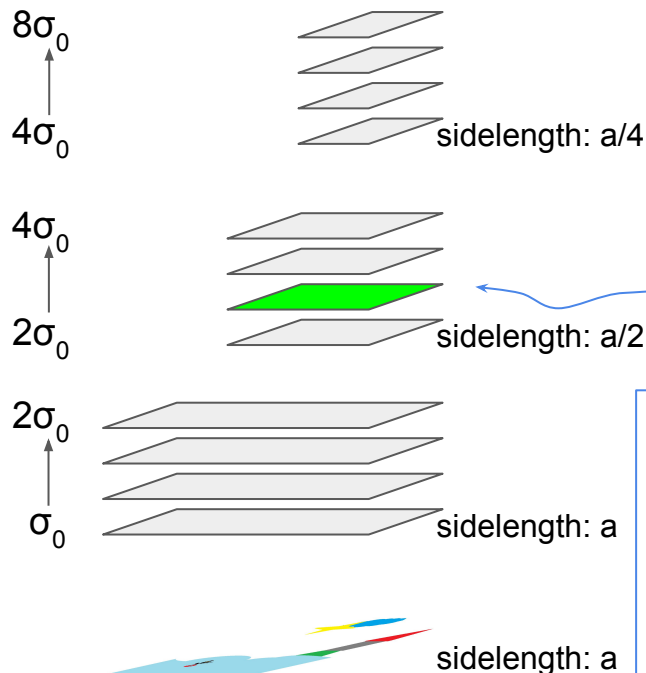
σ : parameter of Gaussian smoothing

of layers (inside an octave) --- *this is the desired/final number of layers*

of octaves


How to produce a pyramid?

the pyramid initially:



inside one octave:

- **LoG*** would be welcomed as edge enhancement on the layers ...
- ... but **DoG**** is computationally far more cheaper



an actual local neighborhood:
they are used for local extrema search
--- see upcoming slide

a layer inside *octave2*, blurred with a Gaussian, using $\hat{\sigma} \in [2\sigma_0, 4\sigma_0]$ standard deviation

$$\hat{\sigma} = \sigma_{scaled} * 2^{\frac{s-2}{NOL}}$$

where

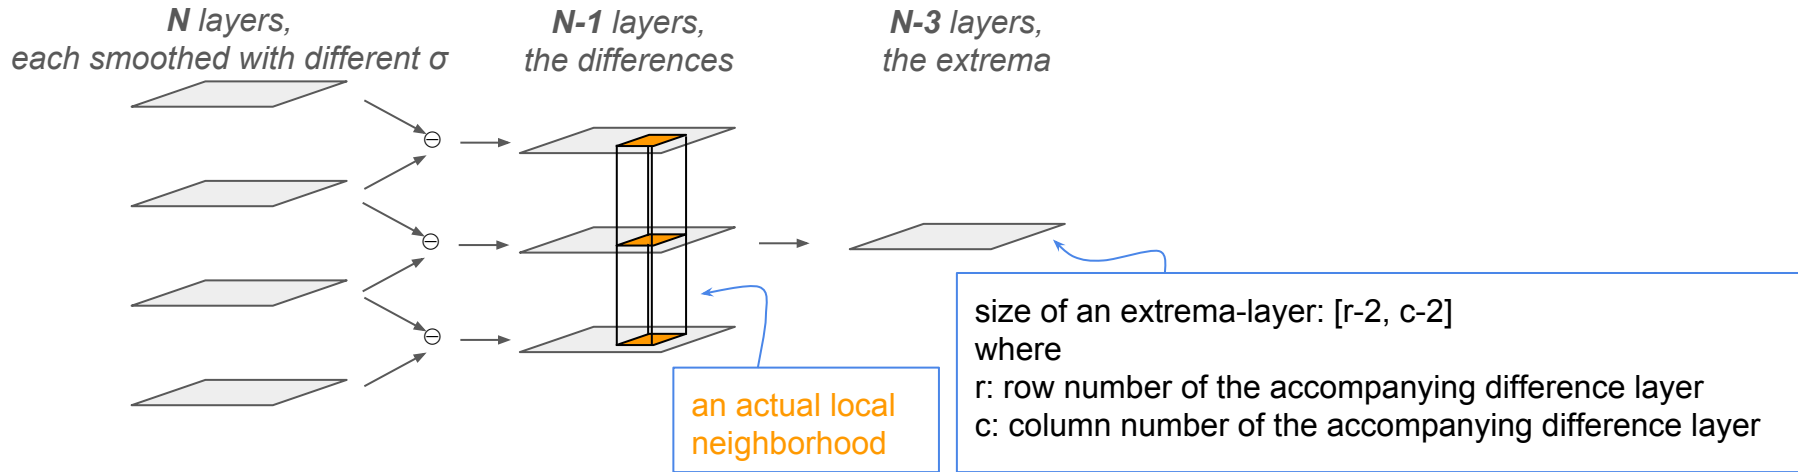
σ_{scaled} is the octave-starting σ ($= 2\sigma_0$)
 s is the layer index inside the octave ($= 2$)
 NOL is the number of layers inside one octave

* LoG: Laplacian of Gaussian
 (see slides 53-56 of Lecture 2)
 ** DoG: difference of Gaussians

Local extrema search

... the magical moment of the foam rubber made from polystyrene

If a central pixel in the 3x3x3 local neighborhood is extremum (min or max), it will be a keypoint-candidate.



Thoughts about the different layer-numbers inside the same octave:

1. initially we use N layers (and smooth them with different σ -s)
2. after calculating the differences, **we lose 1 layer**
3. during the extrema-search: we are interested only in "true" local neighborhoods, so the top and bottom difference layers in the octave could not be analyzed for extrema (as they do not have real upper/lower neighboring layer, resp.), furthermore: for the same reason, the boundary pixels could not be extrema, too \Rightarrow **we lose 2 more layers**

This "**N-3**" will be **equal to** the **desired number of layers** (Definitions on slide 5)

Now please
download the 'Lab 11' code package
from the
[submission system](#)

Exercise 1

Complete the **function** `sift_build_pyramid` in which implement the body of the loop:

- input and output of the function is the `params` structure, which contains all of the necessary parameters in its fields,
- the header of the loop iterates through all of the `octaves` in the pyramid, at every `octave` please:
 - calculate a `scaled_sigma` ($\sigma_{\text{scaled}} = \text{DefaultSigma} * 2^{\text{octave}-1}$) and a `scaled_image` (`InputImage` resized with the built-in `imresize`, at scale $2^{-\text{octave}+1}$),
 - create the 3D array of layers (`layers`) with size `[r, c, NumberOfLayers+3]`, where `r` and `c` are the row- and column-numbers of the `scaled_image` itself,
 - with a `for`-loop (cycle index: `s`), iterate along the 3rd dimension of `layers`, at every step
 - calculate the actual standard deviation (σ^{\wedge} on Slide 6, `NOL = NumberOfLayers`),
 - overwrite the `s`-th layer of `layers` with the gauss-filtered `scaled_image` (use the built-in `imgaussfilt` with standard deviation σ^{\wedge}),
- after the outer `for`-loop, save the layer-wise differences of `layers` into the `octave`-th cell of the `pyramid` field of the `params` struct (use the built-in `diff` on `layers`, with parameters indicating 1st (order) differences and along 3rd dimension --- see Help).

Test 1

Run the **script SIFT_script**:

if the implementation is correct, you should get the struct `params` and an empty vector `SIFT` in the workspace.

Let's examine the different layers inside the pyramid --- all the details are contained in the `params` struct:

```
figure;  
subplot(2, 2, 1); imagesc(params.pyramid{1}(:, :, 1)); title('octave1, layer1');  
subplot(2, 2, 2); imagesc(params.pyramid{1}(:, :, end-1)); title('octave1, layer last-1');  
subplot(2, 2, 3); imagesc(params.pyramid{2}(:, :, 1)); title('octave2, layer1');  
subplot(2, 2, 4); imagesc(params.pyramid{2}(:, :, end-1)); title('octave2, layer last-1');
```

Things to observe:

- upper row: same image-size, different blurring,
- lower-left: half of the side-length, but same blurring with upper-right.

Exercise 2

Complete the **function** `sift_find_extrema` in which implement the body of the loop:

- input and output of the function is the `params` structure, which contains all of the necessary parameters in its fields,
- the already initialized 3D array `map` will contain the places of local extrema: if an element of it has value 1, it means that the corresponding point of the `layers` of the actually processed octave is a local minimum/maximum,
- write 3 nested `for`-loops to visit all **inner** elements of `layers` (1st: through 3rd dimension, 2nd: through rows, 3rd: through columns):
 - extract the 3x3x3 local neighborhood (`nbhd`) of the actual element being visited,
 - reorganize `nbhd` to a column-vector (operator `:`),
 - search for the *index* of the minimum and maximum value, if one of those indices is equal to 14, then the central pixel is an extrema (please consider: the properly reshaped 3x3x3 array will have its central element at the middle of the column-vector) \Rightarrow update `map` if necessary.


NB - along **3rd dim.**: `layers` contains `NumOfLayers`+2 slices while `map` should be indexed from 1 till `NumOfLayers`.

Test 2

Run the **script** `SIFT_script` again:

if the implementation is correct, you should got the struct `params` and the array `SIFT` with size `Px128` in the workspace.

($200 < P < 800$)



Let's examine the different fields of `params`: `extrema`, `filtered` and `eliminated`:

```
figure;  
imshow(params.InputImage);  
hold on;  
spy(params.extrema{1}(:, :, 3), 'g');      % type A  
spy(params.filtered{1}(:, :, 3), 'co');    % type B  
spy(params.eliminated{1}(:, :, 3), 'yx'); % type C
```

Things to observe:

- type A: extrema even on the blue sky,
- type B: low-contrast extrema already filtered out, but edges are still problematic due to aperture problem,
- type C: only complex/2D structures, which can be located accurately.

THE END